

Name: Aadhya

Mail id: aadhyag101@gmail.com

PROFESSIONAL SUMMARY

- Overall 12 years of experience in Architecture/Designing of Azure Data development projects in Defined Benefits Implementations, Big Data, Microsoft Azure Cloud, Python, PySpark, On-prem to Azure cloud solutions.
- Hands on experience in working with Azure Cloud and its components like **Azure Data Factory, Azure Data Lake Gen2, Azure Blob Storage, Azure Databricks, Azure Synapse, Azure Logic Apps, Azure Key Vault, NiFi, Snow pipe, Informatica, Ataccama, Azure Machine Learning**
- Implemented delta tables, delta live tables, autoloader and job API using Databricks.
- Proficient in leveraging Azure Databricks and Spark for distributed data processing and transformation.
- Implemented data pipelines using spark and hands-on experience in spark performance tuning.
- Implemented data pipelines using hive and leveraged various file formats such as ORC, Parquet, Avro and delimited
- Implemented data ingestion using Stream sets, ADF, Sqoop and PySpark .
- Worked on azure data bricks, pyspark, HDInsight, U-SQL, T-SQL, spark SQL, Azure ADW and hive used to load and transform data.
- Hand-on experience in using Azure Data Lake as Source and pulled data using Azure Polybasic.
- Used stored procedure, lookup, execute pipeline, data flow, copy data, azure function features in ADF.
- Experienced in azure Data Lake storage gen2 to store excel files, parquet files and retrieve using data using Blob API.
- Developed and maintained scalable web applications using various technologies, including NoSQL databases like MongoDB and Cassandra.
- Implemented ETL pipelines within and outside of a data warehouse using Python and Snowflakes Snow SQL Querying Snowflake using SQL.
- Development of scripts using Unix, Python, etc. for loading, extracting, and transforming data.
- Performed ETL using Azure Data Bricks. Migrated on-premise Oracle ETL process to Azure Synapse Analytics.
- Assist with production issues in Data Warehouses like reloading data, transformations, and translations.
- Experience in Developing a Database Design and Reporting Design based on Business Intelligence and Reporting requirements.
- Hand-on experience in Developing different APIs to connect with the different Media Data feeds to get the Data using Azure Web Job and Functions integration with Cosmos DB.
- Experience with Snowflake Multi - Cluster Warehouses.
- Experience in using Snowflake Clone and Time Travel.
- Designed and implemented scalable data architecture solutions, including data models, integration strategies, and ETL processes, to optimize data accessibility and support business intelligence initiatives.
- Participates in the development, improvement and maintenance of snowflake database applications.
- Built the Logical and Physical data model for snowflake as per the changes required.
- In-depth knowledge of Snowflake Database, Schema and Table structures.
- Proven track record in performance optimization and capacity planning to ensure scalability and efficiency.
- Proficient in scripting languages such as Python and Scala.
- Hands on working experience and developing large scale data pipelines using spark and hive
- Hands-on experience in setting up workflow using Apache Oozie workflow engine for managing and scheduling Hadoop jobs.
- I am experienced in setting up workflow using Apache Oozie workflow engine for managing and scheduling Hadoop jobs, and importing and exporting the data using SQOOP from HDFS to Relational Database systems.
- I have optimized query performance in Hive using bucketing and partitioning techniques, and have extensive hands-on experience tuning spark Jobs.
- Proficient in utilizing Purview, a unified data governance platform, to establish robust data management practices.
- Skilled in leveraging Purview's data mapping capabilities to create comprehensive data lineage and gain insights into data assets.

- Experienced in implementing Purview for efficient data discovery, cataloging, and classification, ensuring adherence to data governance policies and regulations.
- Proficient in utilizing Purview's data mapping features to identify data relationships, dependencies, and transformations, enhancing data understanding and decision-making
- Skilled in working with Hive, SparkSQL, Kafka, and Spark Streaming for ETL tasks and real-time data processing.
- Have a good working experience in Hadoop, HDFS, Map-Reduce, Hive, TeZ, Python, PySpark.
- Skilled in ensuring data quality and integrity through validation, cleansing, and transformation operations.
- Built and maintain pipelines and run ETL workloads directly on a data lakeEnable real-time point-of-sale analytics
- Highly proficient in Agile methodologies, including JIRA for project management and reporting.
- Implemented robust data security measures ensuring the confidentiality, integrity, and availability of sensitive data, while adhering to industry best practices and compliance standards.
- Designed and implemented scalable data architecture solutions, optimizing data accessibility, integration, and governance for enhanced business intelligence and analytics capabilities.

EDUCATION

- Bachelors in computer science – 2011 -SRM university
- Master in computer science -2012-University of Texas Alrington

TECHNICAL SKILLS

Azure Services	Azure data Factory, Azure Data Bricks, snowflake, Logic Apps, Functional App, Snowflake, Azure DevOps
Hadoop Distribution	Cloudera, Horton Works
Big Data Technologies	MapReduce, Hive, Tez, Python, PySpark, Scala, Kafka, Spark streaming, Oozie, Sqoop, Zookeeper
Languages	SQL, PL/SQL, Python, HiveQL, Scala.
Web Technologies	HTML, CSS, JavaScript, XML, JSP, Restful, SOAP
Operating Systems	Windows (XP/7/8/10), UNIX, LINUX, UBUNTU, CENTOS.
Build Automation tools	Ant, Maven
Version Control	GIT, GitHub.
IDE &Build Tools, Design	Eclipse, Visual Studio.
Databases	MS SQL Server 2016/2014/2012, Azure SQL DB, Azure Synapse. MS Excel, MS Access, Oracle 11g/12c, Cosmos DB

WORK EXPERIENCE

Role: Azure Snowflake data engineer

Duration: Apr 2021 to

Present

Client: PepsiCo, Texas.

Responsibilities:

- Designed and implemented scalable data ingestion pipelines using Azure Data Factory, ingesting data from various sources such as SQL databases, CSV files, and REST APIs.
- Developed data processing workflows using Azure Databricks, leveraging Spark for distributed data processing and transformation tasks.
- Ensured data quality and integrity by performing data validation, cleansing, and transformation operations using Azure Data Factory and Databricks.
- Designed and implemented a cloud-based data warehouse solution using Snowflake on Azure, leveraging its scalability and performance capabilities.
- Created and optimized Snowflake schemas, tables, and views to support efficient data storage and retrieval for analytics and reporting purposes.
- Collaborated with data analysts and business stakeholders to understand their requirements and

implemented appropriate data models and structures in Snowflake.

- Developed and optimized Spark jobs to perform data transformations, aggregations, and machine learning tasks on big data sets.
- Leveraged Azure Synapse Analytics to integrate big data processing and analytics capabilities, enabling seamless data exploration and insights generation.
- Configured event-based triggers and scheduling mechanisms to automate data pipelines and workflows.
- Implemented data lineage and metadata management solutions to track and monitor data flow and transformations.
- Identified and resolved performance bottlenecks in data processing and storage layers, optimizing query execution and reducing data latency.
- Implemented partitioning, indexing, and caching strategies in Snowflake and Azure services to enhance query performance and reduce processing time.
- Conducted performance tuning and capacity planning exercises to ensure the scalability and efficiency of the data infrastructure.
- Implemented data models and optimized queries to improve application performance and reduce response times.
- Leveraged NoSQL database features to handle complex data structures and semi-structured data.
- Worked closely with the development team to design efficient data storage and retrieval strategies for high-traffic applications.
- Hands on programming experience in scripting languages like python and Scala.
- Involved in running all the Hive scripts through Hive on Spark and some through SparkSQL
- Hands on experience in using Kafka, Spark streaming, to process the streaming data in specific use cases.
- Developed a data pipeline using Kafka, Spark, and Hive to ingest, transform and analyze data.
- Designed and implemented real-time data processing solutions using Kafka and Spark Streaming, enabling the ingestion, transformation, and analysis of high-volume streaming data.
- Develop transformation logic using snowpipeline and experience in building snowpipe.
- Developed Spark core and Spark SQL scripts using Scala for faster data processing.
- Deep Knowledge of various NiFi Processors
- Working with JIRA to report on Projects, and creating sub tasks for Development, QA, and Partner validation.

Environment: Azure Databricks, Data Factory, Snowflake, Logic Apps, Functional App, Ataccama, Snowsql, MS SQL, Oracle, HDFS, MapReduce, YARN, Spark, Hive, SQL, Python, Scala, PySpark, shell scripting, JIRA, and kafka.Snowpipe,Nifi, Profisee, Informatica.

Role: Azure Snowflake data engineer
2021

Duration: Dec 2018 to Mar

Client: Wipro, Sunnyvale, California.

Responsibilities:

- Implemented end-to-end data pipelines using Azure Data Factory to extract, transform, and load (ETL) data from diverse sources into Snowflake.
- Designed and implemented data processing workflows using Azure Databricks, leveraging Spark for large-scale data transformations.
- Built scalable and optimized Snowflake schemas, tables, and views to support complex analytics queries and reporting requirements.
- Developed data ingestion pipelines using Azure Event Hubs and Azure Functions to enable real-time data streaming into Snowflake.
- Worked on Snowsql,Snowpipe and created snowpipe for continuous dataload.
- Leveraged Azure Data Lake Storage as a data lake for storing raw and processed data, implementing data partitioning and data retention strategies.
- Utilized Azure Blob Storage for efficient storage and retrieval of data files, implementing compression and encryption techniques to optimize storage costs and data security.
- Integrated Azure Data Factory with Azure Logic Apps for orchestrating complex data workflows and

triggering actions based on specific events.

- Implemented data governance practices and data quality checks using Azure Data Factory and Snowflake, ensuring data accuracy and consistency.
- Implemented data replication and synchronization strategies between Snowflake and other data platforms using Azure Data Factory and Change Data Capture techniques.
- Developed and deployed Azure Functions for data preprocessing, data enrichment, and data validation tasks in data pipelines.
- Implemented advanced analytics and machine learning workflows using Azure Machine Learning and Snowflake, enabling predictive analytics and data-driven insights.
- Designed and implemented data archiving and data retention strategies using Azure Blob Storage and Snowflake's Time Travel feature.
- Developed custom monitoring and alerting solutions using Azure Monitor and Snowflake Query Performance Monitoring (QPM) for proactive identification and resolution of performance issues.
- Integrated Snowflake with Power BI and Azure Analysis Services for creating interactive dashboards and reports, enabling self-service analytics for business users.
- Optimized data pipelines and Spark jobs in Azure Databricks for improved performance, including tuning of Spark configurations, caching, and leveraging data partitioning techniques.
- Implemented data cataloging and data lineage solutions using tools like Azure Purview and Apache Atlas to provide a comprehensive understanding of data assets and their relationships.
- Collaborated with cross-functional teams including data scientists, data analysts, and business stakeholders to understand data requirements and deliver scalable and reliable data solutions.
- Actively involved in staying updated with the latest features and advancements in Purview, ensuring optimal utilization and continuous improvement of data governance processes.
- Successful in implementing Purview's data mapping capabilities to enhance data integration efforts, streamline data workflows, and optimize data-driven initiatives.
- Strong analytical and problem-solving skills, combined with a deep understanding of Purview, to identify and resolve data governance challenges and drive data-driven decision-making.

Environment: Azure Databricks, Data Factory, Logic Apps, Snowflake, Purview data mapping ,Functional App, Snowflake, MS SQL, Oracle, HDFS, MapReduce, YARN, Spark, Purview unified data governance, Hive, SQL, Python, Scala, PySpark, shell scripting, GIT, JIRA, Jenkins, kafka.

Role: Big data Developer

Duration: oct 2017 to Nov 2018

Client: Baker Hughes, Houston, TX.

Responsibilities:

- Prepared an ETL framework using Sqoop, Pig, and Hive to bring in data from various sources and make it available for consumption.
- Processed HDFS data and created external tables using Hive, along with developing scripts for table ingestion and repair for reuse across the project.
- Developed ETL jobs using Spark and Scala to migrate data from Oracle to new MySQL tables.
- Utilized Spark (RDDs, DataFrames, Spark SQL) and Spark-Cassandra Connector APIs for various tasks, including data migration and business report generation.
- Developed a Spark Streaming application for real-time sales analytics.
- Analyzed SQL scripts and designed solutions using PySpark.
- Extracted data from various data sources into HDFS using Sqoop.
- Handled data import from various sources, performed transformations using Hive and MapReduce, and loaded data into HDFS.
- Extracted data from MySQL into HDFS using Sqoop.
- Implemented automation for deployments using YAML scripts for streamlined builds and releases.
- Worked with Apache Hive, Apache Pig, HBase, Apache Spark, Zookeeper, Flume, Kafka, and Sqoop.
- Implemented data classification algorithms using MapReduce design patterns.

- Extensively worked on creating combiners, partitioning, and distributed cache to enhance the performance of MapReduce jobs.
- Utilized Git and GitHub repositories to maintain the source code and enable version control.

Environment: Sqoop, MYSQL, HDFS, Apache Spark Scala, Hive Hadoop, Cloudera, HBASE, Kafka, MapReduce, Zookeeper, Oozie, Data Pipelines, RDBMS, Python, PySpark, shell script, Ambari, JIRA, API.

Role: Data Warehouse Developer

Duration: May 2015 to Feb 2017

Client: HTC Global Svc, Michigan.

Responsibilities:

- Expert understanding of the ETL/ELT and architectural principles of data integration and data warehousing
- Works with on/off shore teams to insure 'best practice' standards and process
- Long hours of sitting and viewing computer monitors
- Responsible for managing detailed Data analysis and translates Business requirements into logical data models
- Manage and participate as a senior Data Analyst gathering of source systems, processing logic, content and operational system usage
- Experience in database management, administration, and development, with a significant percentage working with SQL Server
- Lead project to create reusable transformations, mapplets and mappings in Informatica to create consistency in the application of business rules.
- Worked on project to reduce the number of staging and transformation layers between source data and reporting.
- Physical and Logical Database designing and modeling experience.
- Backup/Restore databases, Monitor, and optimize databases.
- Data Modeling, ETL and data mapping across projects.

Environment: MS SQL Server 2014, SSIS, SSAS, SSRS, SQL Profiler, Power BI, Performance Point Server, MS Office, SharePoint, SQL/PLSQL

Role: Data Warehouse Developer

Duration: Feb 2013 to April 2015

Client: Well care Tampa, FL

Responsibilities

- Experience in developing complex store procedures, efficient triggers, required functions, creating indexes and indexed views for performance.
- Create technical solutions and design for Data Warehouse.
- Define overall ETL architecture including key designs on integration standards such as loading, real time/batch, CDC, data validation, parallelization etc
- Perform analysis and capacity planning to define 'best in class' solutions
- Work with management to align solutions and data integration with business strategy and objectives
- Apply broad in depth business and technical knowledge advance technical direction
- In a direct role as a designer and/or ETL/ELT with Ab Initio and Teradata
- Create complex technical solutions and design utilizing Ab Initio, Teradata and other technologies
- Lead technical discussions impacting solutions and data integration technology
- Experience with SQL Analysis Services, data warehousing, SharePoint and Spotfire desired but not required
- Analytic Data Modeling (Star Schema, Normalized Schema, SCD's, and Lookups)
- ETL and Data Integration skills (Bulk Loads, Temp Tables, and multi-step Transformational SQL)
- Performance Troubleshooting and Optimization (Tuning for Columnar DW-style DB's)
- Python and Scripting Language (DB libraries or dictionaries)
- Develop ETL / ELT solutions in a Data Warehouse environment

Environment: MS SQL Server, Visual Studio, SSIS, SSRS, Informatica, SQL/PLSQL Team Foundation server, Git.