

RISHITHA CHERUKURI

Phone: 913-238-4206

Email: rishithacherukuri61@gmail.com

LinkedIn: linkedin.com/in/rishitha-cherukuri-3877021b9

PROFESSIONAL SUMMARY

Results-driven Data Engineer with over **3 years of real-time consulting experience** designing and deploying scalable data platforms, cloud-based pipelines, and analytics solutions. Proven track record of delivering clean, governed, and production-grade data solutions in **financial services** and **supply chain** domains. Proficient in **Python, PySpark, Apache Airflow, SQL, Snowflake, AWS, and Azure**. Experienced in building real-time data ingestion, automation, and orchestration frameworks that directly impact business intelligence, risk reporting, and operational KPIs. Adept at collaborating with cross-functional teams, gathering data requirements, and driving continuous delivery within Agile environments.

TECHNICAL SKILLS

- **Languages & Scripting:** Python, SQL, Shell Scripting
 - **Big Data Frameworks:** Apache Spark, Hadoop, Hive
 - **Orchestration & Scheduling:** Apache Airflow, Azure Data Factory
 - **Cloud Platforms:** AWS (S3, Lambda, Redshift), Azure (Databricks, Blob Storage)
 - **Data Warehousing & Databases:** Snowflake, PostgreSQL, Oracle, MySQL
 - **DevOps & Tools:** Git, Jenkins, Docker, Kafka, REST APIs
 - **BI & Visualization:** Power BI, Tableau
 - **Modeling & Design:** Star Schema, Snowflake Schema, ERD
 - **Methodologies:** Agile, CI/CD, Data Quality, Data Governance
-

PROFESSIONAL EXPERIENCE

Data Engineer

Client: Fifth Third Bank | Kentwood | MI

May 2024 – Present

Project: Enterprise Financial Data Lake Modernization

- Spearheaded the design and implementation of a cloud-native **data lake** architecture using **AWS S3, Redshift, and Snowflake**, modernizing legacy systems to support real-time analytics.
- Developed more than 50 robust **ETL pipelines** using **PySpark** and **Apache Airflow** to ingest structured, semi-structured, and unstructured data from 20+ upstream systems.
- Migrated 100+ legacy ETL workflows from on-premise SQL jobs to fully automated **Airflow DAGs**, reducing latency by 35% and improving maintainability.
- Built reusable Python libraries for ingestion, validation, and transformation logic, enabling rapid onboarding of new datasets.
- Integrated **Redshift** and **Snowflake** for hybrid warehousing and reporting use cases; optimized storage with clustering keys and compression strategies.
- Created RESTful **data ingestion APIs** to facilitate real-time streaming data from fraud detection and transaction monitoring platforms.
- Developed **data quality rules**, null checks, threshold validations, and duplication checks using custom Python-based frameworks.
- Implemented monitoring and alerting for pipeline health using Airflow sensors, email/SNS notifications, and custom audit logs.

- Collaborated with business analysts and finance teams to define **data dictionaries**, business rules, and lineage for financial reporting.
- Contributed to development of **role-based access policies** and **PII masking** procedures for compliance with internal and external regulations.
- Delivered data marts that serve **Power BI** and **Tableau** dashboards used by risk, compliance, and executive leadership.
- Supported production systems during off-hours and performed root cause analysis for failed loads and job regressions.
- Participated in Agile ceremonies, including sprint planning, daily stand-ups, and retrospectives as a senior consultant.

Environment: Python, PySpark, Apache Airflow, AWS S3, AWS Lambda, Redshift, Snowflake, Tableau, Jenkins, Git, Power BI, JIRA

Data Engineer

Client: Caterpillar | Hyderabad | India

June 2021 – July 2023

Project: Global Supply Chain Analytics Platform

- Contributed to the development of a **supply chain analytics platform** by building streaming and batch pipelines using **Apache Spark**, **Kafka**, and **Python**.
- Created data extraction modules to source real-time inventory and shipment data from internal ERPs and third-party APIs.
- Optimized multiple Hive queries and legacy ETL jobs by introducing **partitioning**, **bucketing**, and caching logic.
- Designed and implemented **Kafka topics and consumers** for capturing data from IoT-enabled warehouse devices and delivery tracking systems.
- Developed data validation scripts in Python to check schema mismatches, field-level anomalies, and timestamp accuracy before ingestion.
- Converted 25+ ETL workflows from manual SQL jobs into **modular Apache Airflow DAGs**, increasing job reliability and observability.
- Built **lineage and compliance documentation**, mapping business fields to raw sources to support audit and regulatory teams.
- Designed Tableau dashboards that visualize supplier performance, lead time trends, and backorder rates for supply chain leadership.
- Performed regular unit tests, code reviews, and pair programming sessions with senior developers to ensure production readiness.
- Maintained version-controlled pipelines via Git and coordinated with DevOps to ensure smooth deployment via Jenkins.
- Collaborated with global data engineers across US and EMEA teams, adhering to global code standards and delivery timelines.

Environment: Python, Apache Spark, Kafka, Hive, Tableau, Oracle, Apache Airflow, Git, SQL, Jenkins

EDUCATION

Master of Science in Computer Science

University of Central Missouri, USA

August 2023 – May 2025

Bachelor of Technology in Computer Science Engineering

Koneru Lakshmaiah University, India

July 2019 – May 2023

CERTIFICATIONS

- Databricks Lakehouse Fundamentals
- SQL for Data Science – Coursera
- Microsoft Certified: Azure Data Engineer Associate (In Progress)
- AWS Certified Data Analytics – Specialty (Planned)