



Satwiki Kaja

Data Engineer



3167307203



satwikikaja20@gmail.com

Results-driven Data Engineer with 5+ years of experience in designing and building robust data pipelines, warehouses, and analytics solutions. Skilled in leveraging **AWS**, **Azure**, and **GCP** to architect scalable data platforms that support business intelligence and real-time analytics.

EDUCATION: Wichita State University, Master's in Data Science, USA, from 2022 - 24

- **PROFILE SUMMARY:** Results-driven **Data Engineer** with over **5+ years of experience** in designing and managing end-to-end **data pipelines**, optimizing data architectures, and supporting enterprise-level analytics initiatives across diverse domains.
- Skilled in **Database Design & Development, Data Modeling, Data Warehousing, Normalization, SQL Reporting, Report Automation, ETL, Agile/Scrum, OLAP/OLTP**, and **KPIs**.
- Proficient in **Python** and **SQL** for **data transformation, processing, and extraction**, with expertise in **ETL tools** and **Macros**.
- Hands-on development using **Python IDEs (PyCharm, Sublime Text, Jupyter Notebook)** and libraries like **Pandas, NumPy, SQLAlchemy, PySpark, Boto3, Matplotlib**.
- Strong experience in building **Spark applications** with **RDD transformations, Spark Core, Spark MLlib, Spark Streaming, and Spark SQL**.
- Expertise in **AWS ecosystem – EC2, S3, EMR**, with experience in **application migration and data migration** to **Snowflake Data Warehouse**.
- Skilled in **SnowSQL, Hadoop/CDH/MapR distributions**, and **Hadoop components: HDFS, Name Node, Data Node, MapReduce**.
- Proficient with **ELK Stack (Elasticsearch, Logstash, Kibana)** for **data analytics, centralized logging, and visualizations**; configured **PagerDuty alerts** for **Elasticsearch, Kafka, Logstash, microservices**.
- Developed **Spark Streaming modules** for ingesting data into **Data Lake** and built robust **ETL packages (SSIS)** for database migration.
- Strong reporting background: created **SSRS reports** (Tabular, Crystal, Matrix, Drill-Down, Cross Tab) and **tabular cubes** for **analytics and BI**.
- Followed **PEP-8 standards**, validated code with **Pylint**, and applied **QA methodologies (SDLC, STLC)** for reliable delivery.
- Experienced in **Test Planning, Test Case Design, Manual & Automated Testing**, with strong knowledge of **Agile/Waterfall lifecycles**.
- Designed and executed automation frameworks using **Selenium (WebDriver, IDE, RC, Grid), TestNG, and Maven**, with scripting in **Java**.
- Conducted diverse testing types: **Functional, Regression, Integration, UAT, Smoke, Sanity, Compatibility, Cross-Browser**.

TECHNICAL SKILLS:

- **Programming & Scripting:** Python, Java, Scala, SQL, Bash, R
- **Big Data Technologies:** Hadoop, Apache Spark, Hive, HDFS, Presto, Pig, Flink
- **ETL/ELT Tools:** Informatica, Talend, Apache NiFi, dbt, SSIS
- **Data Warehousing:** Snowflake, Redshift, Azure Synapse, Big Query, Teradata, Vertica
- **Data Integration & Orchestration:** Apache Airflow, Luigi, AWS Step Functions, Azure Logic Apps, GCP Data Fusion
- **Streaming & Real-Time Processing:** Apache Kafka, AWS Kinesis, Azure Event Hubs, Google Pub/Sub, Spark Streaming

- **Database Systems:** PostgreSQL, MySQL, Oracle, SQL Server, MongoDB, Cassandra, DynamoDB, Cosmos DB, HBase
- **Business Intelligence Tools:** Tableau, Power BI, Looker, AWS QuickSight, Google Data Studio
- **DevOps & CI/CD:** Jenkins, GitLab CI/CD, Azure DevOps, AWS Code Pipeline, Docker, Kubernetes
- **Machine Learning Integration:** AWS SageMaker, Azure ML, GCP Vertex AI
- **Workflow Automation:** Terraform, CloudFormation, Ansible
- **Version Control & Collaboration:** Git, GitHub, GitLab, Bitbucket
- **Cloud Platforms:**
 - **AWS:** Redshift, S3, Glue, Lambda, EMR, Athena, DynamoDB, Kinesis, CloudFormation
 - **Azure:** Azure Data Factory, Synapse Analytics, Azure Storage, Databricks, Cosmos DB, Logic Apps, Azure Functions
 - **Google Cloud (GCP):** Big Query, Dataflow, Pub/Sub, Cloud Storage, Cloud Composer, Cloud Functions

WORK EXPERIENCE:

Dairy Farmers of America

Azure Data Engineer | Kansas City, USA | Jun 2024 – Present

Description: Dairy Farmers of America (DFA) is a leading national dairy cooperative in the dairy industry by marketing and processing raw milk, as well as producing a wide range of dairy products. Designing and developing data integration solutions for cloud data warehouses and data lakes, optimizing data collection and flow across functional teams, and optimizing data processing and managing data storage using Azure services like Azure SQL Database and Azure Data Lake. **Responsibilities:**

- Implementation of **end-to-end data solution** on Azure using **Azure Databricks, ADF, DW and PowerBI**.
- Designed a robust **data modelling environment** using **Databricks on Azure**, enabling consumers to operate highly descriptive **Notebooks** in a fully **governed environment**.
- **Migration** of large **data sets** to **Databricks (Spark)**, create and administer **cluster**, configure **data pipelines**, load data from **ADLS Gen2** to **Databricks** using **ADF pipelines**.
- Hands-on experience writing **Notebooks** in **Databricks** using **Python/Spark SQL** for complex **aggregations, transformations, schema operations**. Proficient with **Databricks Delta** and **DataFrames**.
- Designed and implemented scalable **ETL pipelines** to process multiple **data types** (structured, unstructured) and **file formats (JSON, CSV, text-delimited)**.
- Built **pipelines** to integrate **telemetry streams** and **data sources**, crafting a unified **data model** for **analytics and reporting**.
- Created **temporary views**, curated **data** in **tables**, and performed **Databricks job configuration** with **ETL notebook refactoring**.
- Developed **Azure Data Factory pipelines** to pull data from **multiple sources** and load into **Azure SQL Database** with **triggers** for daily execution.
- Created **SSIS packages** for **ETL**, transforming source data into **dimension/fact tables** and implementing **Slowly Changing Dimensions (SCD)**.
- Migrated **on-prem databases** into **Azure SQL**, handled **Huge Data Transfers** using **DTS, SSIS, BULK INSERT** with **configuration files** and **variables**.
- Designed and loaded **tabular model cubes** on **Azure SQL Database** for **Power BI reports and dashboards**.
- Created/modified **SQL views, stored procedures**, and optimized **ETL scripts** for reporting needs.
- Extensively used **DAX** for complex **calculated measures/columns** in **Power BI** and implemented **row-level security** for role-based access.

Environment: Azure Data Factory, Azure Synapse, SQL Server, Azure Blob Storage, Azure Machine Learning, MLflow, Power BI, Python, SQL, Scala, Apache Spark, Azure Databricks, Kafka, Hadoop, Hive, Cosmos DB, Snowflake, Pandas, NumPy, Python scripting, Azure Event Hubs, Stream Analytics, and Power BI, Azure Monitor.

Parker-Hannifin Corporation

AWS Data Engineer| Cleveland, Ohio, USA| July 2023 – May 2024

Description: Parker-Hannifin Corporation, headquartered in Cleveland, Ohio, is a global leader in motion and control technologies, providing precision-engineered solutions for a wide range of industrial and aerospace markets. Implement data transformation logic using SQL and Python to create derived metrics, calculations, and aggregate data for analysis and design and implement data models in AWS Redshift or other data warehouses to optimize query performance.

Responsibilities:

- Worked with **Ecommerce, Marketing Salesforce, Sales & Ops, Manufacturing, Customer Support** teams to implement projects like **Data Warehouse, Data Engineering, Data Integration Automation, Process Design, API Enablement, Analytics, Data Quality**.
- Solid experience implementing **complex business rules** by creating **re-usable transformations** and robust **mappings** using **Talend transformations (tConvertType, tSortRow, tReplace, tAggregateRow, tUnite)**.
- Involved in **system failure analysis**, identifying **root causes**, and recommending corrective actions.
- Worked on **Hive and BigQuery (BQ)** for **data export, transformation**, and file generation across **analytical formats**.
- Experience in **ETL (Extraction, Transformation, Loading)** from **heterogeneous systems: JSON, XML, Flat Files, Excel, Oracle, MySQL, SQL Server, Salesforce Cloud, API endpoints**.
- Handled **data importing** in **Talend**, performed **transformations** using **Hive, Pig, Spark**, and loaded data into **LUDP**.
- Imported data from **BigQuery** to **LUDP Hive tables**.
- Created **POCs** for **operational logs** across **GCP, AWS, LUDP** using **Cassandra**.
- Partnered with **Ecommerce DevOps engineers** to automate **log filtering** using **AWS S3, Gsutil, Splunk**.
- Collaborated with **Mobile Ecommerce team** to build a **Data Mart** on **AWS Redshift**.
- Developed **POC for Business Intelligence Chatbot** data modeling using **Neo4j, MongoDB, Bigtable**.
- Created **Python scripts** to perform **data cleaning** across multiple tables.
- Created **Hive External Tables**, loaded data, and queried using **HQL**.
- Developed **RESTful APIs** for **Customer Care System** using **Talend ESB** to provide seamless access to **customer and product data**.
- Responsible for **Data Modeling** and development of **Business Intelligence Chatbot** using **Python Flask** and **Google Cloud**, enabling real-time access to **business KPIs**.
- Improved **daily job performance** via **data cleaning, query optimization, table partitioning**.
- Created **automation processes** for **Inventory & Sales Data Distribution** using **Talend + BigQuery**, sending activation reports.
- Designed **data ingestion process** using **BigQuery & GCS** to process **200 GB+ daily manufacturing test data**.
- Built **API extraction process** for **third-party vendor orders & distributor data**.
- Automated **email attachment extraction** and **data push** into **BigQuery**.
- Performed **data mapping** (source-to-target) and converted **JSON → XML (Accord Format)** using **Talend Data Mapper + TXMLMap**.
- Created **execution plans** in **TAC**.
- Developed **Talend Quality Checks (Joblets)** based on business rules.
- Built **Talend Mappings** to populate **Dimension & Fact Tables**.
- Developed **file movement jobs** to vendor servers (**monthly, weekly, daily**) using **Talend**.
- Analyzed **customer data** to run **Push Notification Campaigns**, improving adoption rates by **28%** using **Tableau**.
- Identified defective **manufacturing stations** and test processes using **Smart Factory methods**, reducing costs by **31%**.

Environment: GCP, Big Query, ETL workflows, Power BI, Cloud SQL, CI/CD, Terraform, Cloud Build, GitOps, Hadoop, Hive, Spark, AWS S3, AWS Redshift, AWS Glue, RDBMS, Sqoop, Impala, Cloud Logging, Cloud Monitoring.

JPMorgan Chase

GCP Data Engineer | Mumbai, India | Nov 2020 – June 2022

Description: JPMorgan Chase & Co. is one of the largest and most well-known financial institutions, offering a broad range of services in investment banking, financial services. Created and maintained efficient and scalable data pipelines using GCP services like [Big Query](#), [Dataflow](#), [Cloud Storage](#), and [Pub/Sub](#), and implemented data quality checks, monitoring, and validation processes to ensure data accuracy and consistency.

Responsibilities:

- **Engineered and optimized ETL workflows** integrating structured and semi-structured data from **Cloud Storage**, **Cloud SQL**, and external financial feeds to power **credit risk models** and **trading analytics**.
- **Developed dynamic Power BI dashboards** connected to **BigQuery** datasets, enabling **real-time visualization** of trading activity, customer insights, and **financial KPIs** for executive reporting.
- **Automated CI/CD pipelines** and **infrastructure provisioning** using **Terraform**, **Cloud Build**, and **GitOps** principles, ensuring compliance with **security and audit standards**.
- Applied **Python** and **SQL** for **data wrangling**, **feature engineering**, and **analytics development** supporting **fraud detection**, **customer segmentation**, and **real-time decision-making** in high-volume banking systems.
- Migrated **legacy data warehouse systems** to **scalable big data platforms** by integrating **Hadoop**, **Hive**, and **Spark**, ensuring compatibility with modern **cloud ecosystems**.
- Enabled **multi-cloud data orchestration** by integrating **AWS S3**, **AWS Redshift**, and **AWS Glue** with **GCP processes**, enhancing **cross-platform data accessibility** and **analytics efficiency**.
- Built **custom ingestion and transformation pipelines** for **risk model inputs** and **regulatory filings**, ensuring **data lineage and governance** with **Data Catalog**, **IAM roles**, and **service accounts**.
- Utilized **Sqoop** for **RDBMS-to-Hadoop transfers**, leveraged **Impala** for **fast SQL queries**, and managed distributed services with **Zookeeper** to support **regulatory reporting analytics**.
- Ensured **audit-ready environments** by implementing **data validation**, **access control**, and **real-time monitoring** through **Cloud Logging**, **Cloud Monitoring**, and custom alerting scripts.

Environment: GCP, BigQuery, ETL workflows, Power BI, Cloud SQL, CI/CD, Terraform, Cloud Build, GitOps, Hadoop, Hive, Spark, AWS S3, AWS Redshift, AWS Glue, RDBMS, Sqoop, Impala, Cloud Logging, Cloud Monitoring

Glenmark Pharmaceuticals Ltd | Data Engineer | Mumbai, India | June 2019 – Oct 2020

Description: Glenmark Pharmaceuticals Ltd is a research-driven multinational pharmaceutical company with innovative drug discovery alongside complex generics and high-value specialty medicines. Developing and maintaining databases and data systems and working with medical experts to develop new policies and improve processes and identifying promising drug candidates, and optimizing the drug development process.

Responsibilities:

- Designed and developed robust data pipelines using **Apache Spark** and **Apache Kafka** to ingest, transform, and integrate clinical research and pharmaceutical sales data from various internal lab systems and external APIs into a centralized data lake on **AWS S3**, enabling real-time analytics and regulatory reporting.
- Implemented scalable **data warehousing solutions** using **Amazon Redshift** and **AWS Glue**, enabling cross-functional teams to query, derive insights from high-volume datasets related to drug formulation, distribution, and R&D operations.

- Utilized **SQL**, **Python**, and **Shell scripting** to clean, validate, and enrich large datasets from drug discovery platforms and commercial operations, ensuring high data accuracy and consistency for pharmacovigilance and compliance tracking.
- Built and maintained interactive dashboards and performance reports using **Power BI**, **Looker**, and **Tableau** to deliver real-time visibility into trial success metrics, production output, and market trends for high-value specialty medicines.
- Engineered **Hadoop-based data pipelines** using **HDFS**, **Hive**, and **MapReduce** for processing and analysing unstructured research data, including compound screening results and clinical study outcomes, and scalable data handling.
- Integrated **Zookeeper** with **Kafka** to coordinate distributed systems for consistent pharmaceutical data streaming and ensured fault-tolerant processing of manufacturing and clinical trial feeds across global teams.
- Automated data workflows using **Apache Airflow** and **AWS Lambda** to orchestrate ETL pipelines across drug discovery, development, and commercial datasets, reducing manual intervention and increasing pipeline reliability.
- Optimized **SQL queries** and **Redshift clusters** to enhance the performance of business intelligence workloads, reducing query execution time and enabling fast insights into therapeutic segment performance and inventory forecasting.

Environment: Apache Spark, Apache Kafka, AWS S3, Amazon Redshift, AWS Glue, SQL, Python, Shell scripting, Power BI, Looker, Tableau, HDFS, Hive, MapReduce, Zookeeper, Kafka, Apache Airflow, AWS Lambda, SQL queries