



Professional Summary:

- Results-driven Azure Data Engineer with 7 years of experience in data engineering and analytics, including 5 years of expertise in Azure services and 2 years in Tableau development.
- Designed, developed, and maintained scalable data pipelines and architectures using **Azure Data Factory, Azure Synapse, Databricks**, and related Azure services.
- I specialized in designing scalable data Lakehouse architecture, orchestrating dynamic data pipelines, and implementing end-to-end CI/CD automation.
- Defined data models, schemas, and partitioning strategies for performance optimization.
- Integrated on-premises and cloud data sources.
- Proven success in migrating enterprise data from on-prem to Azure cloud platforms, developing real-time ingestion and transformation workflows using ADF, Azure Databricks, Delta Live Tables, and Synapse Analytics.
- I am skilled in data modeling, governance with Unity Catalog & Purview, and integrating APIs from external systems. Microsoft Certified (DP-900, DP-700) and passionate about solving complex business problems with reliable, performant data solutions.
- Built streaming pipelines with **Azure Event Hubs, IoT Hub**, or **Azure Stream Analytics**.
- Automate data workflows and infrastructure provisioning using **Azure DevOps, Terraform**, or **ARM/Bicep templates**.
- Adept at translating business needs into technical solutions and collaborating across cross-functional teams. Experienced in managing end-to-end data pipelines, ensuring data quality, governance, and security compliance in hybrid environments.
- Used **Databricks, Azure Synapse Spark**, or **HDInsight** for big data transformations.
- Optimized storage performance tiers and compute resources.
- Ensured data security, privacy, and compliance by configuring **RBAC**, encryption, and network security controls.
- Documented data flows, architecture designs, and operational procedures.

Certifications:

- ☐ Microsoft Certified: Fabric Data Engineer Associate (DP-700)
- ☐ Microsoft Certified: Azure Data Fundamentals (DP-900)

Education:

- Master's Degree, KL University, India.

Technical Skills:

Cloud & Big Data:	Azure Data Factory, Azure Databricks, Delta Live Tables, Synapse Dedicated SQL Pools, ADLS Gen2, Azure Event Hub, Azure Purview, Azure API Management
Programming & Scripting:	Python, PySpark, SQL, T-SQL, REST APIs
Data Integration:	Azure DevOps, ARM Templates, Git, CI/CD
Architecture:	Medallion Architecture (Bronze, Silver, Gold), Star Schema, Real-Time Streaming
ETL & Tools:	ADF Pipelines, Synapse Pipelines, Dataflows Gen2, Cosmos DB, NoSQL

Visualization: Tableau (Desktop & Server), Power BI (basic)
Data Governance: Unity Catalog, Azure Purview, Data Quality (DQ) Rules, Metadata Management
Scheduling & Monitoring: ADF triggers, Tumbling Window, Event-based triggers, Databricks Job Clusters, Logging with Azure Monitor

Professional Experience:

Client: CVS

Duration: Oct'2023 – Till date

Role: Azure Data Engineer

Project: Enterprise Master Data Management (MDM) Pipeline

Project Description: Designed and implemented a scalable, enterprise-grade Master Data Management (MDM) pipeline to consolidate, cleanse, and unify customer records from multiple heterogeneous sources, enabling a single source of truth for analytics and downstream applications. The solution leveraged the Medallion Architecture (Bronze, Silver, Gold) on Azure Data Lake, with orchestration via Azure Data Factory, transformation into Azure Databricks (PySpark), and storage in Azure SQL Database for BI and reporting consumption. The project ensured data quality, governance compliance, and optimized query performance for analytical workloads.

Responsibilities:

- Developed **dynamic, parameterized Azure Data Factory (ADF) pipelines** to ingest datasets from SQL Server, SFTP, and Blob Storage into **ADLS Gen2 Bronze Layer**, enabling a **60% reduction in manual data preparation time**.
- Implemented **event-based & scheduled triggers** for automated ingestion workflows.
- Built **PySpark notebooks** in Azure Databricks to cleanse, normalize, and standardize attributes in the **Silver Layer**.
- Applied **fuzzy matching UDFs** (fuzzywuzzy) to detect near-duplicates and improve matching accuracy.
- Used **window functions** (row_number() over Window.partitionBy...) to retain the latest, trusted records.
- Defined **MDM survivorship rules** prioritizing source trust hierarchy & last update timestamps.
- Loaded Gold Layer tables into **Azure SQL Database** using JDBC connectors from Databricks, reducing **Power BI report load times by 40%** through performance tuning.
- Loaded Gold Layer tables into **Azure SQL Database** using JDBC connectors from Databricks.
- Designed and implemented **T-SQL performance tuning** strategies (indexes, partitioning) to support **millions of records with minimal performance degradation**.
- Created **role-based access control (RBAC)** in Azure SQL with read-only roles for BI consumers, ensuring **compliance with GDPR and internal data privacy standards**.
- Version-controlled ADF pipelines, Databricks notebooks & SQL scripts in **Azure DevOps Git**, ensuring **traceability and controlled deployments**.

Environment: Cloud Platform: Microsoft Azure, Ingestion: Azure Data Factory (ADF) – Linked Services, Datasets, Mapping Data Flows, Storage: Azure Data Lake Storage Gen2 (Bronze/Silver/Gold), Processing: Azure Databricks (PySpark, Delta Lake, Fuzzy Matching, Window Functions), Serving: Azure SQL Database (T-SQL, Indexing, Partitioning, RBAC), Version Control / CI/CD: Azure DevOps Git Repos, YAML Pipelines, Security: Azure SQL RBAC, Managed Identities, Azure AD Integration, Analytics Consumption: Power BI (DirectQuery / Import modes).

Client: Belk, Charlotte, NC

Duration: June'2021 – Sept'2023

Role: Azure Data Engineer

Project: Real-Time Salesforce Integration & Streaming Analytics with Synapse

Responsibilities:

- Developed robust ADF pipelines to ingest Salesforce CRM data using REST API integration with dynamic parameters and lookup configurations.
- Mapping Data Flows used to flatten and normalize complex nested JSON structures received from Salesforce endpoints.
- Configured Linked Services, datasets, and parameterized pipeline templates to promote reusability and scalability across environments.
- Implemented pipeline scheduling using Tumbling Window triggers and configured alerting, retries, and logging with Azure Monitor integration.
- Leveraged Azure Databricks workflows to clean, deduplicate, and enrich the raw data using PySpark and stored the output in Delta format within ADLS Gen2.
- Built Delta Live Tables in Databricks to manage CDC (Change Data Capture) operations and maintain versioned data.
- Modeled curated data in star schema within Synapse Dedicated SQL Pools, building dimension and fact tables for analytical workloads.
- Worked closely with business users to define KPIs and published cleaned datasets to Power BI for reporting and dashboard development.
- Applied role-based access controls using Unity Catalog and documented data flows, dependencies, and scheduling in operational runbooks.
- Enabled incremental data loads and recovery mechanisms, reducing pipeline failure rates and improving SLAs across the reporting platform.
- Improved pipeline reliability by 95% through automated retries and alerting, meeting critical SLAs.
- Reduced manual Salesforce data preparation efforts by 70% through reusable parameterized pipelines.

Environment: Microsoft Azure, Azure Data Factory (ADF) – Linked Services, Datasets, REST API integration, Mapping Data Flow, Storage: Azure Data Lake Storage Gen2 (Bronze/Silver/Gold), Processing: Azure Databricks (PySpark, Delta Live Tables, CDC handling), Modeling & Serving: Synapse Dedicated SQL Pools (Star Schema, Fact & Dimension tables), Version Control / CI/CD: Azure DevOps Git Repos, YAML Pipelines, Security: Unity Catalog (RBAC), Azure AD Integration, Analytics Consumption: Power BI (DirectQuery & Import modes).

Client: Target Corporation, Minneapolis, MN

Duration: Dec'2019 – May'2021

Role: Azure Data Engineer

Project: Enterprise Data Lakehouse Migration to Azure (ADF + Databricks + Synapse)

Responsibilities:

- Migrated different data sources from on-prem SQL and SFTP into ADLS Gen2 using dynamic and parameterized ADF pipelines with triggers.
- Automated transformations using Azure Databricks workflows with PySpark scripts to build silver and gold layer Delta tables.
- Implemented end-to-end CI/CD pipeline using Azure DevOps and YAML for deploying ADF, Key Vault, and Databricks assets.
- Built analytical views using Synapse Dedicated SQL Pools and served curated data to Tableau and Power BI.
- Applied Unity Catalog for access control and Purview for lineage, metadata classification, and data cataloging.

- Developed stored procedures for incremental refresh of target tables in the gold layer and scheduled them using Azure Data Factory triggers to refresh data at regular intervals.
- Integrated Azure Functions to execute stored procedures as part of the Logic Apps workflow execution.
- Designed and implemented data quality (DQ) queries, scheduled them for periodic execution, and configured alert mechanisms to notify the team in case of any outliers.
- Optimized query performance in Azure Synapse Analytics by distributing smaller tables across different nodes to minimize data shuffling.
- Enabled near-time analytics and reduced manual efforts by 85% with robust monitoring and retry logic.
- Created HLD (High-Level Design) and LLD (Low-Level Design) documentation for ADF workflows and Databricks jobs.
- Configured alerting and retry policies for long-running and failure-prone pipelines using ADF fault tolerance settings.

Client: Citi Group, NYC, NY

Duration: Aug'2017 – Oct'2019

Role: Tableau Developer

Project: Tableau Dashboards for Executive Sales Reporting

Responsibilities:

- Designed 15+ Tableau dashboards using data from Synapse and Azure SQL for sales and operations teams.
- Created calculated KPIs, dynamic filters, and parameter-based views with RLS (Row-Level Security).
- Partnered with stakeholders to understand business metrics and iteratively delivered insights.
- Conducted performance tuning and publishing to Tableau Server with access control.
- Improved visibility into regional and product-level sales performance with drill-down insights.

You can contact this candidate at:
a8f3a70d-4bbf-42e6-8479-fa9e2a227671@contact.monster.com