

ZAID SIDDIQUI

Alpharetta, GA

zaidstddiqui1223@gmail.com | (678) 698-3859 | <https://www.linkedin.com/in/zaid-siddiqui-631343176/>

SENIOR DATA ENGINEER | CLOUD & BI SOLUTIONS ARCHITECT

Results-driven Data Engineer with 9+ years of experience architecting scalable data platforms, cloud-native pipelines, and BI solutions across Azure, AWS, and GCP. Adept at building robust ETL/ELT frameworks, optimizing performance, and delivering actionable analytics with tools like ADF, Databricks, DBT, and Snowflake. Proven track record of reducing latency, improving pipeline reliability, and integrating data governance. Collaborative team player and mentor who thrives in fast-paced, cross-functional environments.

Core Competencies: Azure Data Factory, Databricks, Snowflake, DBT, Python (PySpark, Pandas), Power BI, Airflow, Azure DevOps, Azure Purview, Fabric, Redshift, Synapse, SQL Tuning, CI/CD, Data Governance, Medallion Architecture

TECHNICAL SKILLS

- **Cloud Platforms:** Azure (ADF, Databricks, Synapse, Purview, Fabric), AWS (Glue, Redshift, Lambda, S3), GCP (BigQuery, Dataflow)
 - **Data Engineering:** Snowflake, DBT, Airflow, Airbyte, Kafka, Delta Lake, SSIS, Azure Monitor
 - **Programming:** Python (Pandas, PySpark, Boto3), SQL (Tuning, Procedures), C# (SSIS Scripts)
 - **BI & Reporting:** Power BI (DAX, RLS), Tableau, SSRS, DOMO
 - **DevOps:** Azure DevOps, GitHub, CI/CD Pipelines, Terraform, ARM Templates
 - **Other:** Star/Snowflake Schemas, REST APIs, Data Modeling, Parquet, Spark Optimization, HIPAA Compliance, RBAC
-

PROFESSIONAL EXPERIENCE

KPMG — *Senior Data Engineer*

New York, NY | Sep 2023 – Present

- Architected hybrid data pipelines across Azure and AWS using ADF, Databricks, and AWS Glue, ingesting data from APIs, S3, SQL, and RDS into Synapse and Redshift.

- Designed Microsoft Fabric Lakehouse pipelines and Spark Notebooks for large-scale batch and streaming workloads.
- Integrated Delta Lake and Apache Airflow for advanced orchestration.
- Developed modular DBT models in Snowflake with automated testing and data quality checks.
- Created Power BI dashboards for sales and supply chain KPIs using DAX and RLS.
- Automated CI/CD using Azure DevOps and GitHub; deployed infrastructure with ARM templates.
- Implemented data governance via Azure Purview and RBAC; established monitoring via Azure Monitor.
- Tuned Spark jobs and SQL queries, achieving sub-second response times and a 40% reduction in ETL latency.
- Improved pipeline reliability through test-driven development in DBT workflows.

Comcast — *BI Data Engineer*
Alpharetta, GA | Aug 2019 – Sep 2023

- Engineered scalable ADF pipelines handling over 1TB daily from SQL Server to Synapse.
- Built and optimized PySpark ETL processes in Databricks with automated deployment via Azure DevOps.
- Designed and deployed semantic models (SSAS Tabular) and dashboards using Power BI and Tableau.
- Implemented Medallion architecture in Azure Data Lake (bronze/silver/gold layers).
- Tuned Snowflake and Redshift data warehouses for performance and cost efficiency.
- Reduced data processing time by 30% by reengineering ETL transformations.
- Led migration of legacy ETL systems to modern Azure-based architecture.

UPS — *SQL Developer*
Alpharetta, GA | Jun 2016 – Jul 2019

- Built ETL pipelines using SSIS, ADF, and Databricks for logistics and operational reporting.
 - Designed SSAS multidimensional and tabular models with complex DAX and MDX logic.
 - Developed SSRS and Tableau reports with dynamic drill-down capabilities.
 - Automated data pipelines and Power BI report refreshes, cutting manual workload by 30%.
-

PROJECT HIGHLIGHTS

- **Global Sales Dashboard (Power BI):** Integrated Synapse and Snowflake to deliver real-time insights with dynamic filtering and secure RLS.
 - **Data Lakehouse Modernization:** Built Microsoft Fabric pipelines integrating S3, Delta Tables, and Fabric Lakehouse for a Fortune 500 client.
 - **Data Governance Implementation:** Applied Azure Purview for full data lineage, cataloging, and role-based access control.
 - **API Ingestion Framework:** Developed reusable Python framework for ingesting and transforming API data into Snowflake using DBT.
 - **Real-Time Inventory Analytics:** Leveraged PySpark, Kafka, and Event Hubs to enable sub-minute tracking of inventory events, reducing latency by 50%.
-