

KRISHNA GUTTA

Apex, NC | Mobile: 8722651579 | Email: krigutta@gmail.com

PROFESSIONAL SUMMARY

- Results-driven Data Engineer with 8+ years of experience designing and implementing scalable data architectures, batch and real-time processing solutions, and enterprise-grade ETL workflows across cloud and on-premises environments.
- Strong hands-on experience with Informatica PowerCenter, SnapLogic, NiFi, and SQL Server Integration Services (SSIS) for developing high-performance ETL pipelines, integrating diverse data sources, and ensuring data integrity.
- Proficient in a wide range of technologies including AWS (S3, Redshift, Lambda, EMR, Glue, CodeBuild), Azure (Databricks, Data Factory, Synapse, Data Lake), Snowflake, Databricks, Apache Airflow, Spark, Hive, Pig, Python, R, SQL, and Apache Kafka.
- Extensive expertise in migrating legacy ETL pipelines (Oracle, SQL Server) to cloud-native platforms, including API-based and JSON transformations for loading data into SQL and cloud data warehouses, improving processing time, scalability, and system reliability.
- Experienced in developing robust, automated data pipelines and workflows for both batch and streaming data, leveraging orchestration tools and frameworks such as Airflow, NiFi, Talend, and Ab Initio.
- Successfully integrated Kafka and Lambda for real-time and event-driven architectures, supporting low-latency, high-throughput streaming data applications, and API-to-SQL integrations.
- Strong skills in SQL (T-SQL, PL/SQL), Oracle, SQL Server, and Redshift for querying, modeling, performance tuning, and optimizing analytical workloads.
- Built and managed large-scale data lakes, data warehouses, and BI solutions, using tools like Hive, Pig, Spark, Snowflake, Power BI, Tableau, and Amazon QuickSight to support analytics and machine learning use cases.
- Adept at enforcing data engineering policies, data governance practices (RBAC), and quality standards to maintain accuracy, reliability, and compliance at scale.
- Experienced in ML model development for anomaly detection and predictive analytics using Python-based tools and cloud platforms.
- Strong collaboration and communication skills; worked closely with business stakeholders to gather requirements, translate objectives into technical solutions, and align initiatives with organizational goals.
- Flexible to support on-call production issues and incident response, ensuring high availability and performance of mission-critical pipelines.

EDUCATION

Master of Science in **Management information Systems**

Aug 2022 – May 2024

Northern Illinois university, Dekalb, IL

SKILLS

- **Programming & Scripting:** Python, SQL, PL/SQL, Shell scripting, R
- **Data Engineering & ETL Tools:** Informatica PowerCenter, SnapLogic, Apache NiFi, AWS Glue, SQL Server Integration Services (SSIS)
- **Cloud Platforms & Services:** AWS (EMR, Redshift, Lambda, S3, CodeBuild, CloudFormation, EventBridge, Kinesis, SNS, SQS, Sagemaker, CloudWatch, CodePipeline), Azure (Databricks, Data Lake, Synapse, Blob Storage), Databricks
- **Big Data & Processing Frameworks:** Apache Spark, PySpark, Hive, Pig, Kafka (producers/consumers), Airflow
- **Data Warehousing & Storage:** Snowflake, Amazon Redshift, SQL Server, Oracle, Hadoop HDFS
- **Containerization & Orchestration:** Docker, Kubernetes
- **APIs & Data Formats:** REST APIs, JSON transformations, API-to-SQL integration
- **Visualization & Reporting:** Amazon QuickSight, Power BI, Azure Databricks visualizations
- **Version Control & Collaboration:** GitHub, Jira, Confluence
- **Data Governance & Best Practices:** RBAC, data quality assurance, stakeholder collaboration

WORK EXPERIENCE

Data Engineer

Jun 2024-Present

Client: Costco

- Designed and implemented data-intensive applications in the cloud using AWS services such as CodeBuild, Glue, Lambda, Step Functions, S3, CloudFormation, Databricks, and SDKs, resulting in improved data processing efficiency and scalability.
- Worked on migration of ETL processes from Oracle SQL to AWS, leveraging Amazon EMR, Databricks, and Redshift for optimized performance, which reduced data processing time and improved system reliability.
- Designed, developed, and maintained ETL workflows using Informatica PowerCenter, Databricks, and SQL Server Integration Services (SSIS), enabling efficient extraction, transformation, and loading of large-scale enterprise data across multiple sources and targets.
- Built and maintained event-driven ETL pipelines using AWS Lambda, S3, Python, JSON parsing, EventBridge, Kinesis, SNS, and SQS, including formatting and loading API-based data feeds into SQL Server and cloud warehouses, automating ingestion and transformation with high fault tolerance.
- Implemented ML-driven data enrichment workflows by integrating Amazon SageMaker into pipelines, enabling scalable model training and real-time inference for analytics and recommendation use cases.
- Integrated Apache Kafka for real-time data ingestion, enabling streaming pipelines that improved data freshness and reduced latency in downstream analytics.
- Containerized data workflows with Docker and deployed scalable microservices on Kubernetes clusters, improving portability, fault tolerance, and deployment consistency across environments.

- Developed Kafka producers and consumers for processing event-driven data streams, ensuring reliable delivery and transformation of high-throughput messages in a decoupled architecture.
- Constructed end-to-end data pipelines with NiFi on EC2, streamlining data flow across systems and integrating external APIs where required.
- Orchestrated and automated complex batch workflows using Apache Airflow, scheduling and monitoring dependencies across AWS and on-premise systems, improving pipeline reliability and reducing manual interventions.
- Utilized Snowflake as a cloud data warehouse to store and analyze large datasets and support business reporting needs.
- Optimized Snowflake virtual warehouse sizing, partitioning, and clustering strategies to reduce compute costs while improving query performance for analytics workloads.
- Developed secure data sharing solutions and implemented Snowflake RBAC policies to maintain compliance and control access across teams.
- Developed DBT models and incremental transformations in Snowflake, streamlining analytics-ready datasets and enabling reproducible, modular data transformations.
- Managed version control with GitHub and tracked project progress using Jira and Confluence, maintaining clear and detailed documentation and collaborating with business stakeholders to gather requirements.

Data Engineer

Jan 2024-May 2024

Client: Macy's

- Migrated and optimized data pipelines to AWS Cloud, leveraging S3, Glue, Lambda, EMR, Databricks, Redshift, Athena, CloudFormation, and SDKs to deliver high-performance, scalable ETL workflows.
- Built real-time and batch data ingestion pipelines using Apache Kafka, AWS Kinesis, EventBridge, SNS, and SQS, enabling low-latency processing of IoT and application data.
- Designed and orchestrated ETL workflows using AWS Glue, SSIS, PySpark, and Databricks, processing structured and semi-structured data (JSON, Parquet) to support enterprise analytics.
- Leveraged Docker for containerized ETL jobs and orchestrated deployments with Kubernetes, ensuring seamless scaling and efficient resource utilization across dev, test, and prod environments.
- Constructed end-to-end ETL and ingestion pipelines using NiFi on EC2, streamlining data flow across systems and integrating external APIs where required.
- Integrated Amazon SageMaker with AWS pipelines to operationalize machine learning models, enabling automated training, batch inference, and real-time predictions for analytics and personalization use cases.
- Engineered serverless solutions with AWS Lambda, SNS, and SQS to enable event-driven data processing and asynchronous message handling.
- Developed ELT workflows in Snowflake to transform raw ingested data into analytics-ready datasets, supporting BI and operational dashboards.
- Integrated Snowflake with AWS-based pipelines to centralize enterprise reporting data, enabling advanced analytics and cross-team collaboration.

- Automated infrastructure provisioning and environment management using Terraform and AWS CDK, improving deployment consistency and reducing setup time.
- Wrote optimized SQL and PLSQL scripts, stored procedures, and views for complex ETL transformations and analytics workloads, improving query performance by 30–50%.
- Managed source control and CI/CD workflows using GitHub, GitLab, CodeBuild, and CodePipeline, enforcing test-driven development (TDD) and improving release quality.

Data Engineer

Aug 2020 – May 2022

Shan Softech Solutions, India

- Extracted and ingested data from various sources into Hadoop Data Lake by creating ETL pipelines using Pig and Hive, which streamlined data processing and improved data accessibility.
- Experienced in transferring data from different data sources into HDFS systems using Kafka producers, consumers, and Kafka brokers, including integration of JSON payloads from APIs.
- Migrated data from Hadoop clusters to cloud platforms using AWS S3, EMR, Elastic Cache, and EC2, which improved data accessibility and reduced storage costs.
- Developed Hive and Pig scripts to perform data transformations, event joins, and filtering, which optimized data processing before storing in HDFS.
- Wrote Python scripts and basic R programs to access databases, handle data wrangling tasks, and execute scripts and commands.
- Contributed to early adoption of Snowflake by building proof-of-concept pipelines to load curated Hadoop datasets into Snowflake for analytics, providing exposure to cloud data warehouse capabilities.
- Responsible for defining the data flow within Hadoop ecosystem and exported the result set from Hive to MySQL using Shell scripts.

Junior Data Engineer

Jul 2017 – Jun 2020

Origin IT Solutions, India

- Developed and maintained ETL data pipelines for automated extraction, transformation, and loading using SQL Server, Azure Databricks, and SSIS to optimize data flow for data analysis.
- Cleaned, preprocessed, and transformed raw data using Python (Pandas, NumPy), R (basic data analysis), and Azure Databricks notebooks, ensuring high data quality and integrity for advanced analytics.
- Built predictive models using machine learning algorithms (e.g., decision trees, regression models) with Python, R (basic), and Azure Databricks to forecast trends and predict future behaviors, enhancing data-driven decision-making.
- Leveraged Azure Databricks to process large datasets, reducing processing times and enabling faster insights generation for business-critical decisions.
- Provided actionable insights and recommendations to senior leadership to support business strategies, optimize operations, and inform key decision-making through data analysis and predictive modeling.